

Prediction Assignment Writeup

Juan Tenopala

19 de diciembre de 2017

Prediction Assignment Writeup

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here:

<http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>

(<http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>) (see the section on the Weight Lifting Exercise Dataset).

Objective

To use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants, to predict the manner in which they did the exercise.

Data

The data for this project come from this source:

<http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>

(<http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>). If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment. The training data for this project are available here: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>) The test data are available here: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)

Introduction

One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, the goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants.

The model

The outcome variable is classe, a factor variable with 5 levels. For this data set, participants were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in 5 different fashions: Class A corresponds to the specified execution of the exercise, while the other 4 classes correspond to common mistakes. Prediction evaluations will be based on maximizing the accuracy and minimizing the out-of-sample error. All other available variables after cleaning will be used for prediction. In this case I will use a Random Forest algorithm because the features of the algorithm are important for classification and prediction.

Code and Results

Packages, Libraries and seed

Load packages and libraries

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.4.2
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
##     margin
```

```
library(rpart)  
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 3.4.2
```

```
set.seed(1234)
```

Load the data sets

Load the data sets and clean them to eliminate NA's

```
trainingset <- read.csv("pml-training.csv", na.strings=c("NA", "#DIV/0!", ""))
testingset <- read.csv("pml-testing.csv", na.strings=c("NA", "#DIV/0!", ""))
dim(trainingset)
```

```
## [1] 19622 160
```

```
dim(testingset)
```

```
## [1] 20 160
```

Preprocessing

Delete columns with all "NA" and the variables that are irrelevant to the analysis.

```
trainingset <- trainingset[, colSums(is.na(trainingset)) == 0]
testingset <- testingset[, colSums(is.na(testingset)) == 0]
trainingset <- trainingset[, -c(1:7)]
testingset <- testingset[, -c(1:7)]
```

Exploratory data analysis

Perform an exploratory data analysis and create a plot to have the first look to the preprocessed data

```
subsamples <- createDataPartition(y=trainingset$classe, p=0.75, list=FALSE)
subTraining <- trainingset[subsamples, ]
subTesting <- trainingset[-subsamples, ]
```

```
plot(subTraining$classe, col="blue", main="5 levels in variable classe of the subTraining
data set", xlab="classe levels", ylab="Frequency")
```

5 levels in variable classe of the subTraining data set



From the graph above, we can see that each level frequency is within the same order of magnitude of each other. Level A is the most frequent with more than 4000 occurrences while level D is the least frequent with about 2500 occurrence

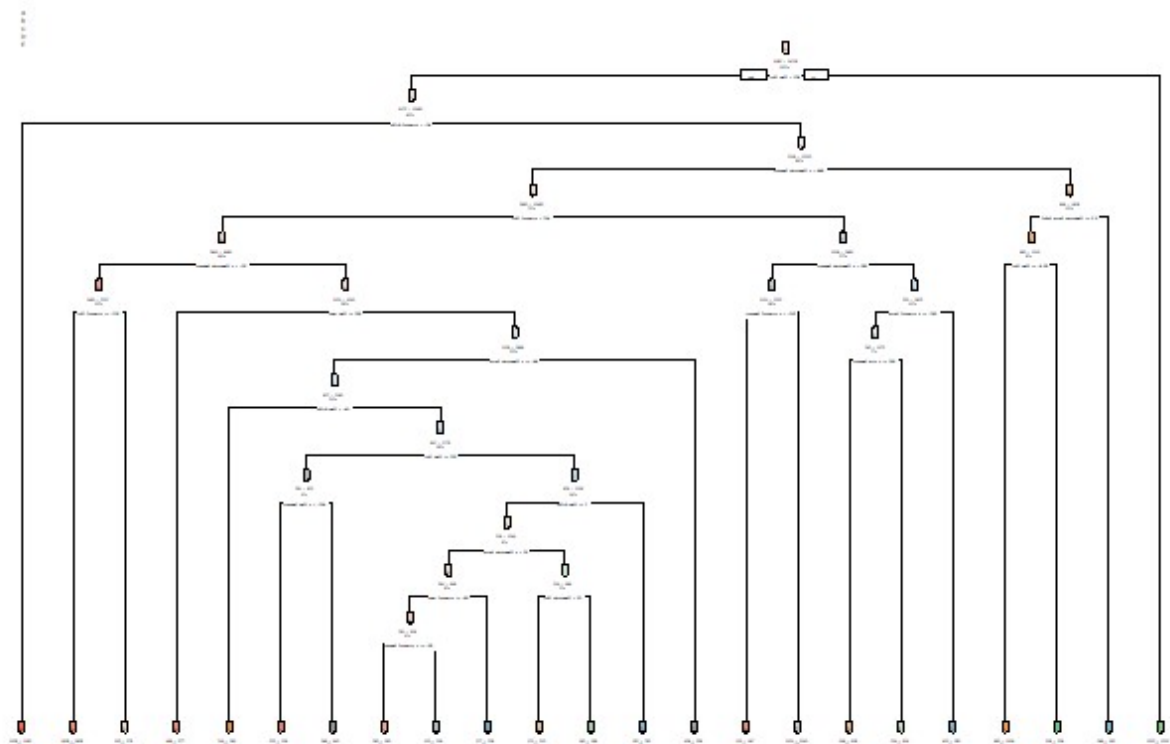
Cross validation

Cross-validation will be performed by subsampling our training data set randomly without replacement into 2 subsamples: subTraining data (75% of the original Training data set) and subTesting data (25%). Our models will be fitted on the subTraining data set, and tested on the subTesting data. Once the most accurate model is chosen, it will be tested on the original Testing data set.

Decision Tree Prediction model

```
DTM <- rpart(classe ~ ., data=subTraining, method="class")
DTP <- predict(DTM, subTesting, type = "class")
rpart.plot(DTM, main="Classification Tree", extra=102, under=TRUE, faclen=0)
```

Classification Tree



```
confusionMatrix(DTP, subTesting$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1235  157   16   50   20
##           B   55  568   73   80  102
##           C   44  125  690  118  116
##           D   41   64   50  508   38
##           E   20   35   26   48  625
##
## Overall Statistics
##
##           Accuracy : 0.7394
##           95% CI : (0.7269, 0.7516)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6697
##           McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8853   0.5985   0.8070   0.6318   0.6937
## Specificity      0.9307   0.9216   0.9005   0.9529   0.9678
## Pos Pred Value   0.8356   0.6469   0.6313   0.7247   0.8289
## Neg Pred Value   0.9533   0.9054   0.9567   0.9296   0.9335
## Prevalence       0.2845   0.1935   0.1743   0.1639   0.1837
## Detection Rate   0.2518   0.1158   0.1407   0.1036   0.1274
## Detection Prevalence 0.3014   0.1790   0.2229   0.1429   0.1538
## Balanced Accuracy 0.9080   0.7601   0.8537   0.7924   0.8307
```

Random Forest Prediction model

```
RFM <- randomForest(classe ~. , data=subTraining, method="class")
RFP <- predict(RFM, subTesting, type = "class")
confusionMatrix(RFP, subTesting$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1394     3     0     0     0
##           B    1   944    10     0     0
##           C    0    2   843     6     0
##           D    0    0    2   798     0
##           E    0    0    0    0   901
##
## Overall Statistics
##
##           Accuracy : 0.9951
##           95% CI : (0.9927, 0.9969)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9938
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9993  0.9947  0.9860  0.9925  1.0000
## Specificity      0.9991  0.9972  0.9980  0.9995  1.0000
## Pos Pred Value   0.9979  0.9885  0.9906  0.9975  1.0000
## Neg Pred Value   0.9997  0.9987  0.9970  0.9985  1.0000
## Prevalence       0.2845  0.1935  0.1743  0.1639  0.1837
## Detection Rate   0.2843  0.1925  0.1719  0.1627  0.1837
## Detection Prevalence 0.2849  0.1947  0.1735  0.1631  0.1837
## Balanced Accuracy 0.9992  0.9960  0.9920  0.9960  1.0000
```

Conclusion

As expected, Random Forest algorithm performed better than Decision Trees. Accuracy for Decision Tree model was: 0.739 (95% CI: (0.7269, 0.7516)) compared to Random Forest model: 0.9951 (95% CI : (0.9927, 0.9969)). The random Forest model is choosen. The accuracy of the model is 0.995. The expected out-of-sample error is estimated at 0.005, or 0.5%. The expected out-of-sample error is calculated as 1 - accuracy for predictions made against the cross-validation set. Our Test data set comprises 20 cases. With an accuracy above 99% on our cross-validation data, we can expect that very few, or none, of the test samples will be missclassified.

```
prediction <- predict(RFM, testingset, type="class")
prediction
```

```
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E
```